

Loan_Analysis

October 17, 2018

```
In [3]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
In [4]: #This dataset contains loan characteristics of a bank's customers.
#The goal is to clean the dataset, create a model that predicts loan outcome(paid or c
#find the characteristics that are highly correlated with safe loan applicanti
```

```
In [5]: #Creating the dataset
data = pd.read_csv('LoansTrainingSet.csv')
```

```
/Users/charliecarrera/anaconda3/lib/python3.6/site-packages/IPython/core/interactiveshell.py:2
interactivity=interactivity, compiler=compiler, result=result)
```

```
In [6]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 256984 entries, 0 to 256983
Data columns (total 19 columns):
Loan ID                256984 non-null object
Customer ID            256984 non-null object
Loan Status            256984 non-null object
Current Loan Amount    256984 non-null int64
Term                  256984 non-null object
Credit Score          195308 non-null float64
Years in current job   245508 non-null object
Home Ownership         256984 non-null object
Annual Income          195308 non-null float64
Purpose               256984 non-null object
Monthly Debt          256984 non-null object
Years of Credit History 256984 non-null float64
Months since last delinquent 116601 non-null float64
Number of Open Accounts 256984 non-null int64
Number of Credit Problems 256984 non-null int64
Current Credit Balance 256984 non-null int64
Maximum Open Credit    256984 non-null object
Bankruptcies           256455 non-null float64
```

```

Tax Liens                256961 non-null float64
dtypes: float64(6), int64(4), object(9)
memory usage: 37.3+ MB

```

```
In [7]: data.head()
```

```

Out[7]:
      Loan ID      Customer ID \
0  000025bb-5694-4cff-b17d-192b1a98ba44  5ebc8bb1-5eb9-4404-b11b-a6eebc401a19
1  00002c49-3a29-4bd4-8f67-c8f8fbc1048c  927b388d-2e01-423f-a8dc-f7e42d668f46
2  00002d89-27f3-409b-aa76-90834f359a65  defce609-c631-447d-aad6-1270615e89c4
3  00005222-b4d8-45a4-ad8c-186057e24233  070bcecb-aae7-4485-a26a-e0403e7bb6c5
4  0000757f-a121-41ed-b17b-162e76647c1f  dde79588-12f0-4811-bab0-e2b07f633fcd

      Loan Status  Current Loan Amount      Term  Credit Score \
0  Fully Paid      11520  Short Term      741.0
1  Fully Paid      3441  Short Term      734.0
2  Fully Paid      21029  Short Term      747.0
3  Fully Paid      18743  Short Term      747.0
4  Fully Paid      11731  Short Term      746.0

      Years in current job  Home Ownership  Annual Income      Purpose \
0      10+ years  Home Mortgage      33694.0  Debt Consolidation
1      4 years  Home Mortgage      42269.0      other
2      10+ years  Home Mortgage      90126.0  Debt Consolidation
3      10+ years      Own Home      38072.0  Debt Consolidation
4      4 years      Rent      50025.0  Debt Consolidation

      Monthly Debt  Years of Credit History  Months since last delinquent \
0      $584.03      12.3      41.0
1      $1,106.04      26.3      NaN
2      $1,321.85      28.8      NaN
3      $751.92      26.2      NaN
4      $355.18      11.5      NaN

      Number of Open Accounts  Number of Credit Problems  Current Credit Balance \
0      10      0      6760
1      17      0      6262
2      5      0      20967
3      9      0      22529
4      12      0      17391

      Maximum Open Credit  Bankruptcies  Tax Liens
0      16056      0.0      0.0
1      19149      0.0      0.0
2      28335      0.0      0.0
3      43915      0.0      0.0
4      37081      0.0      0.0

```

```

In [8]: #Get rid of non-integer units in series
        for ch in ['$ ', ',']:
            data['Monthly Debt'] = [i.replace(ch, '') for i in data['Monthly Debt']]

In [9]: #Make Monthly Debt numeric through list comprehension
        data['Monthly Debt'] = [float(i) for i in data['Monthly Debt']]

In [10]: data['Monthly Debt'].head()

Out[10]: 0      584.03
         1     1106.04
         2     1321.85
         3      751.92
         4      355.18
         Name: Monthly Debt, dtype: float64

In [11]: data['Loan Status'].unique()

Out[11]: array(['Fully Paid', 'Charged Off'], dtype=object)

In [12]: #Change 'Loan Status' unique values to integers, e.g. create dummy variable
        data['Loan Status'] = data['Loan Status'].replace('Fully Paid', 1).replace('Charged Off', 0)

In [13]: data['Loan Status'].head()

Out[13]: 0      1
         1      1
         2      1
         3      1
         4      1
         Name: Loan Status, dtype: int64

In [14]: #The columns 'Term', 'Home Ownership', and 'Purpose' are all categorical
         #Here we create dummy variables and merge them with the dataset
        data2 = pd.merge(data, (pd.get_dummies(data['Term'], drop_first = True)), left_index = True, right_index = True)

In [15]: data3 = pd.merge(data2, (pd.get_dummies(data['Home Ownership'], drop_first = True)), left_index = True, right_index = True)

In [16]: data4 = pd.merge(data3, (pd.get_dummies(data['Purpose'], drop_first = True)), left_index = True, right_index = True)

In [17]: del data4['Term']
         del data4['Home Ownership']
         del data4['Purpose']

In [19]: #Next we deal with missing values for 'Bankruptcies' and 'Tax Liens'
         #Correlation map revealed that both variables are correlated with 'Number of Credit Problems'
         #We fill missing values depending on how many credit problems that individual has had
        data4['Bankruptcies'] = data4['Bankruptcies'].fillna(data4['Bankruptcies'].mean())

In [20]: #Knowing how many credit problems someone has can help us predict how many bankruptcies
        data4.groupby('Number of Credit Problems')['Bankruptcies'].mean()

```

Out[20]: Number of Credit Problems

| | |
|----|----------|
| 0 | 0.000000 |
| 1 | 0.822141 |
| 2 | 0.900234 |
| 3 | 1.039191 |
| 4 | 1.058182 |
| 5 | 1.008000 |
| 6 | 1.214286 |
| 7 | 1.187500 |
| 8 | 0.916667 |
| 9 | 0.100000 |
| 10 | 0.666667 |
| 11 | 0.000000 |

Name: Bankruptcies, dtype: float64

In [21]: *#Filling in missing 'Bankruptcies' values by the number of their credit problems*

```
for i in data4[data4['Bankruptcies'].isnull()].index:
```

```
    data4['Bankruptcies'][i] = data4['Bankruptcies'][data4['Number of Credit Problems']
```

/Users/charliecarrera/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:3: Deprecati
.ix is deprecated. Please use
.loc for label based indexing or
.iloc for positional indexing

See the documentation here:

<http://pandas.pydata.org/pandas-docs/stable/indexing.html#ix-indexer-is-deprecated>

This is separate from the ipykernel package so we can avoid doing imports until

/Users/charliecarrera/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:3: SettingWi

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

This is separate from the ipykernel package so we can avoid doing imports until

In [22]: *#Filling in missing 'Tax Liens' values by the number of their credit problems*

```
for i in data4[data4['Tax Liens'].isnull()].index:
```

```
    data4['Tax Liens'][i] = data4['Tax Liens'][data4['Number of Credit Problems']] == c
```

/Users/charliecarrera/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:3: Deprecati
.ix is deprecated. Please use
.loc for label based indexing or
.iloc for positional indexing

See the documentation here:

<http://pandas.pydata.org/pandas-docs/stable/indexing.html#ix-indexer-is-deprecated>

This is separate from the ipykernel package so we can avoid doing imports until

/Users/charliecarrera/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:3: SettingWi

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>
This is separate from the ipykernel package so we can avoid doing imports until

```
In [23]: #Each Loan is suppose to be unique, therefore duplicates are dropped  
data5 = data4.drop_duplicates(['Loan ID'], keep = 'last')
```

```
In [24]: #Max credit score an individual can have is 800 based off of this particular credit score  
#It appears an extra zero was added to these credit scores  
data5['Credit Score'][data5['Credit Score'] > 800].head(10)
```

```
Out[24]: 341      6600.0  
        349      6760.0  
        420      7460.0  
        522      7320.0  
        623      7270.0  
        846      6690.0  
        926      7230.0  
       1240      7380.0  
       1306      7440.0  
       1317      7390.0  
Name: Credit Score, dtype: float64
```

```
In [25]: #We divide these values by ten to get rid of the extra zero  
#This is done through list comprehension  
data5['Credit Score'][data5['Credit Score'] > 800] = [i/10 for i in data5['Credit Score'] if i > 800]
```

/Users/charliecarrera/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>
This is separate from the ipykernel package so we can avoid doing imports until
/Users/charliecarrera/anaconda3/lib/python3.6/site-packages/pandas/core/generic.py:7620: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>
self._update_inplace(new_data)
/Users/charliecarrera/anaconda3/lib/python3.6/site-packages/IPython/core/interactiveshell.py:2345: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>
exec(code_obj, self.user_global_ns, self.user_ns)

```
In [27]: #There are now no more values above 800  
#Distribution of credit score  
data5['Credit Score'].describe()
```

```
Out[27]: count      167406.000000  
        mean         723.028022
```

```

std          26.648780
min          585.000000
25%          713.000000
50%          732.000000
75%          742.000000
max          751.000000
Name: Credit Score, dtype: float64

```

```

In [28]: #Some individuals had extremely high loan amounts yet low income
#Delete outliers in 'Current Loan Amount'
data6 = data5.drop(data5[data5['Current Loan Amount'] > 1000000].index)

```

```

In [29]: #Next 'Years in current job' is made purely numeric
data6['Years in current job'].head()

```

```

Out[29]: 0    10+ years
1         4 years
2    10+ years
3    10+ years
4         4 years
Name: Years in current job, dtype: object

```

```

In [30]: data6['Years in current job'] = [str(i) for i in data6['Years in current job']]

```

```

In [31]: h = [i.split(' ')[0] for i in data6['Years in current job']]

```

```

In [32]: h = [i.replace('<', '.5') for i in h]
h = [i.replace('n/a', '0') for i in h]
h = [i.replace('10+', '10') for i in h]
h = [float(i) for i in h]

```

```

In [33]: data6['Years in current job'] = h

```

```

In [34]: hh = list(data6['Maximum Open Credit'])

```

```

In [140]: #Now 'Maximum Open Credit is cleaned by getting rid of strings
hh = [str(i).replace('#VALUE!', '0') for i in hh]
hh = [float(i) for i in hh]
data6['Maximum Open Credit'] = hh

```

```

In [141]: data6['Maximum Open Credit'] = [float(i) for i in data6['Maximum Open Credit']]

```

```

In [36]: #Everytime 'Annual Income is null, so is 'Credit Score'
#These null rows are deleted
data[data['Annual Income'].isnull()].head(3)

```

```

Out[36]:
Loan ID \
7    0000afa6-8902-4f8f-b870-25a8fdad0aeb
8    00011dfc-31c1-4178-932a-fbeb3f341efb
12   00029f9f-0cc5-4d4e-aabc-ea4a7fe74e12

```

| | Customer ID | Loan Status | Current Loan Amount | \ |
|----|--------------------------------------|-------------|---------------------|---|
| 7 | e49c1a82-a0f7-45e8-9f46-2f75c43f9fbc | 0 | 24613 | |
| 8 | ef6e098c-6c83-4752-8d00-ff793e476b8c | 1 | 10036 | |
| 12 | afbc2fa3-3bad-4d48-b691-829aed78bad5 | 0 | 17980 | |

| | Term | Credit Score | Years in current job | Home Ownership | \ |
|----|------------|--------------|----------------------|----------------|---|
| 7 | Long Term | NaN | 6 years | Rent | |
| 8 | Short Term | NaN | 5 years | Rent | |
| 12 | Short Term | NaN | < 1 year | Own Home | |

| | Annual Income | Purpose | Monthly Debt | Years of Credit History | \ |
|----|---------------|--------------------|--------------|-------------------------|---|
| 7 | NaN | Business Loan | 542.29 | 17.6 | |
| 8 | NaN | Debt Consolidation | 386.36 | 17.7 | |
| 12 | NaN | Debt Consolidation | 597.50 | 9.9 | |

| | Months since last delinquent | Number of Open Accounts | \ |
|----|------------------------------|-------------------------|---|
| 7 | 73.0 | 7 | |
| 8 | NaN | 7 | |
| 12 | 43.0 | 7 | |

| | Number of Credit Problems | Current Credit Balance | Maximum Open Credit | \ |
|----|---------------------------|------------------------|---------------------|---|
| 7 | 0 | 14123 | 16954 | |
| 8 | 0 | 11970 | 16579 | |
| 12 | 0 | 6817 | 22800 | |

| | Bankruptcies | Tax Liens |
|----|--------------|-----------|
| 7 | 0.0 | 0.0 |
| 8 | 0.0 | 0.0 |
| 12 | 0.0 | 0.0 |

```

In [192]: data7 = data6
          del data7['Months since last delinquent']
          data8 = data7.drop(data7[data7['Annual Income'].isnull()].index)

In [194]: del data8['Loan ID']
          del data8['Customer ID']

In [196]: #Predictions are created for 'Loan Status' aka whether or not the loan was paid off
          #'Loan Status' is categorical, calls for a classification algorithm
          from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
          from sklearn.ensemble import GradientBoostingClassifier
          gbc = GradientBoostingClassifier()

In [ ]: gbc.fit(data8.drop('Loan Status', axis = 1), data8['Loan Status'])

In [770]: #75% accuracy score means
          accuracy_score(data8['Loan Status'], gbc.predict(data8.drop('Loan Status', axis = 1)))

```

```
Out[770]: 0.75739054131743777
```

```
In [779]: gbc.feature_importances_
```

```
Out[779]: array([ 0.10438967,  0.18943464,  0.04284989,  0.18065884,  0.10127577,  
                  0.0714165 ,  0.03295017,  0.00468396,  0.07069356,  0.09599793,  
                  0.01642792,  0.00726124,  0.01304626,  0.01581257,  0.00846362,  
                  0.          ,  0.02838365,  0.          ,  0.0027543 ,  0.00769752,  
                  0.00249725,  0.00094808,  0.          ,  0.00115457,  0.          ,  
                  0.0012021 ])
```