

Real_Estate

October 17, 2018

```
In [3]: #Kaggle competition for building the best model to predict/estimate the sale prices of  
        #Dataset contains 80 characteristics such as location, square_footage, utilities, ect  
        #Beginning showcases exploratory work
```

```
In [ ]: import pandas as pd  
        import numpy as np  
        import matplotlib.pyplot as plt
```

```
In [354]: sample = pd.read_csv('sample_submission.csv')  
          sample.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1459 entries, 0 to 1458  
Data columns (total 2 columns):  
Id          1459 non-null int64  
SalePrice   1459 non-null float64  
dtypes: float64(1), int64(1)  
memory usage: 22.9 KB
```

```
In [4]: train = pd.read_csv('train_kaggle.csv')
```

```
In [5]: test = pd.read_csv('test_kaggle.csv')
```

```
In [6]: data = train.append(test)
```

/Users/charliecarrera/anaconda3/lib/python3.6/site-packages/pandas/core/frame.py:6201: FutureWarning: of pandas will change to not sort by default.

To accept the future behavior, pass 'sort=True'.

To retain the current behavior and silence the warning, pass sort=False

```
sort=sort)
```

```
In [358]: 1459*2
```

```
Out[358]: 2918
```

```

In [359]: 0_type = data['BldgType'].dtype

In [7]: train.groupby('EnclosedPorch')['SalePrice'].mean().head()

Out[7]: EnclosedPorch
0      186856.88099
19     220000.00000
20     122000.00000
24      82000.00000
30     104000.00000
Name: SalePrice, dtype: float64

In [8]: object_columns = [column for column in data.columns if data[column].dtype == data['BldgType'].dtype]

In [9]: data_object = pd.DataFrame()

In [10]: for name in object_columns:
          data_object[name] = data[name]

In [11]: data_object.describe()

Out[11]:
      Alley BldgType BsmtCond BsmtExposure BsmtFinType1 BsmtFinType2 \
count      198      2919      2837          2837          2840          2839
unique        2         5         4            4            6            6
top      Grv1      1Fam         TA           No          Unf          Unf
freq       120      2425      2606          1904           851          2493

      BsmtQual CentralAir Condition1 Condition2 ... MiscFeature \
count      2838      2919      2919      2919 ...           105
unique        4         2         9         8 ...            4
top          TA         Y      Norm      Norm ...          Shed
freq       1283      2723      2511      2889 ...           95

      Neighborhood PavedDrive PoolQC RoofMatl RoofStyle SaleCondition \
count           2919      2919      10      2919      2919          2919
unique           25         3         3         8         6            6
top           NAmes         Y      Ex  CompShg      Gable          Normal
freq           443      2641         4      2876      2310          2402

      SaleType Street Utilities
count      2918      2919      2917
unique        9         2         2
top          WD   Pave   AllPub
freq       2525      2907      2916

[4 rows x 43 columns]

In [365]: for column in object_columns:
          data[column] = data[column].replace(data[column].unique(), range(len(data[column].unique())))

```

```

In [366]: train['GarageType'].unique()

Out[366]: array(['Attchd', 'Detchd', 'BuiltIn', 'CarPort', nan, 'Basment', '2Types'],
              dtype=object)

In [367]: range(len(train['GarageType'].unique()))

Out[367]: range(0, 7)

In [368]: data['GarageType'].unique()

Out[368]: array([0, 1, 2, 3, 4, 5, 6])

In [370]: del data['Alley']
          del data['PoolQC']

In [371]: data.describe()

Out[371]:
```

	1stFlrSF	2ndFlrSF	3SsnPorch	BedroomAbvGr	BldgType	\
count	2919.000000	2919.000000	2919.000000	2919.000000	2919.000000	
mean	1159.581706	336.483727	2.602261	2.860226	0.460774	
std	392.362079	428.701456	25.188169	0.822693	1.088487	
min	334.000000	0.000000	0.000000	0.000000	0.000000	
25%	876.000000	0.000000	0.000000	2.000000	0.000000	
50%	1082.000000	0.000000	0.000000	3.000000	0.000000	
75%	1387.500000	704.000000	0.000000	3.000000	0.000000	
max	5095.000000	2065.000000	508.000000	8.000000	4.000000	

	BsmtCond	BsmtExposure	BsmtFinSF1	BsmtFinSF2	BsmtFinType1	\
count	2919.000000	2919.000000	2918.000000	2918.000000	2919.000000	
mean	0.211716	0.800274	441.423235	49.582248	1.846523	
std	0.676432	1.233050	455.610826	169.205611	1.686056	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	368.500000	0.000000	2.000000	
75%	0.000000	2.000000	733.000000	0.000000	3.000000	
max	4.000000	4.000000	5644.000000	1526.000000	6.000000	

	...	SaleType	ScreenPorch	Street	TotRmsAbvGrd	\
count	...	2919.000000	2919.000000	2919.000000	2919.000000	
mean	...	0.251799	16.062350	0.004111	6.451524	
std	...	0.856761	56.184365	0.063996	1.569379	
min	...	0.000000	0.000000	0.000000	2.000000	
25%	...	0.000000	0.000000	0.000000	5.000000	
50%	...	0.000000	0.000000	0.000000	6.000000	
75%	...	0.000000	0.000000	0.000000	7.000000	
max	...	9.000000	576.000000	1.000000	15.000000	

	TotalBsmtSF	Utilities	WoodDeckSF	YearBuilt	YearRemodAdd	\
--	-------------	-----------	------------	-----------	--------------	---

count	2918.000000	2919.000000	2919.000000	2919.000000	2919.000000
mean	1051.777587	0.001713	93.709832	1971.312778	1984.264474
std	440.766258	0.055510	126.526589	30.291442	20.894344
min	0.000000	0.000000	0.000000	1872.000000	1950.000000
25%	793.000000	0.000000	0.000000	1953.500000	1965.000000
50%	989.500000	0.000000	0.000000	1973.000000	1993.000000
75%	1302.000000	0.000000	168.000000	2001.000000	2004.000000
max	6110.000000	2.000000	1424.000000	2010.000000	2010.000000

	YrSold
count	2919.000000
mean	2007.792737
std	1.314964
min	2006.000000
25%	2007.000000
50%	2008.000000
75%	2009.000000
max	2010.000000

[8 rows x 79 columns]

```
In [372]: data['GarageYrBlt'][data['GarageYrBlt'] > 2018]
```

```
Out[372]: 1132      2207.0
          Name: GarageYrBlt, dtype: float64
```

```
In [373]: #Date entries above the current year follow a pattern
          #Make outlier equal to 2007
          data['GarageYrBlt'][data['GarageYrBlt'] > 2018] = 2007
```

/Users/charliecarrera/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:3: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

This is separate from the ipykernel package so we can avoid doing imports until

```
In [374]: #Change Year variables, 0 is newest year possible, higher number = older house
          k = ['YrSold', 'GarageYrBlt', 'YearBuilt', 'YearRemodAdd']
          for column in k:
              data[column] = [(2018 - i) for i in data[column]]
```

```
In [375]: data[k].describe()
```

```
Out[375]:
```

	YrSold	GarageYrBlt	YearBuilt	YearRemodAdd
count	2919.000000	2760.000000	2919.000000	2919.000000
mean	10.207263	39.959058	46.687222	33.735526
std	1.314964	25.206206	30.291442	20.894344
min	8.000000	8.000000	8.000000	8.000000

25%	9.000000	16.000000	17.000000	14.000000
50%	10.000000	39.000000	45.000000	25.000000
75%	11.000000	58.000000	64.500000	53.000000
max	12.000000	123.000000	146.000000	68.000000

```
In [376]: np.std(data['1stFlrSF'])
```

```
Out[376]: 392.2948646055113
```

```
In [12]: data1 = data
```

```
In [378]: #Normalize:
          #x-min/xmax-xmin
          #0 to 1 scale
```

```
In [379]: is_null = [column for column in data1.columns if data1[column].isnull().sum() > 0]
```

```
In [380]: is_null_less = [column for column in data1[is_null].columns if data1[column].isnull().sum() < 10]
```

```
In [381]: data1[is_null].info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2919 entries, 0 to 1458
Data columns (total 12 columns):
BsmtFinSF1      2918 non-null float64
BsmtFinSF2      2918 non-null float64
BsmtFullBath    2917 non-null float64
BsmtHalfBath    2917 non-null float64
BsmtUnfSF       2918 non-null float64
GarageArea      2918 non-null float64
GarageCars      2918 non-null float64
GarageYrBlt     2760 non-null float64
LotFrontage     2433 non-null float64
MasVnrArea      2896 non-null float64
SalePrice       1460 non-null float64
TotalBsmtSF     2918 non-null float64
dtypes: float64(12)
memory usage: 296.5 KB
```

```
In [382]: data1['BsmtFinSF1'].mean()
```

```
Out[382]: 441.4232350925291
```

```
In [ ]: #Next we fill in missing values
        #Through mean grouping tactics
```

```
In [384]: for column in data1[is_null_less].columns:
          data1[column][data1[column].isnull()] = data1[column].mean()
```

```
/Users/charliecarrera/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

```
In [386]: new_null = [column for column in data1.columns if data1[column].isnull().sum() > 0]
```

```
In [387]: data1[new_null].info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2919 entries, 0 to 1458
Data columns (total 4 columns):
GarageYrBlt      2760 non-null float64
LotFrontage      2433 non-null float64
MasVnrArea       2896 non-null float64
SalePrice        1460 non-null float64
dtypes: float64(4)
memory usage: 114.0 KB
```

```
In [388]: data1['GarageYrBlt'][data1['GarageYrBlt'].isnull()] = data1['YearBuilt'][data1['GarageYrBlt'].isnull()]
```

```
/Users/charliecarrera/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

```
"""Entry point for launching an IPython kernel.
```

```
In [389]: null_now = [column for column in data1.columns if data1[column].isnull().sum() > 0]
```

```
In [390]: data1[null_now].info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2919 entries, 0 to 1458
Data columns (total 3 columns):
LotFrontage      2433 non-null float64
MasVnrArea       2896 non-null float64
SalePrice        1460 non-null float64
dtypes: float64(3)
memory usage: 171.2 KB
```

```
In [392]: data1.ix[5]['MasVnrType']
```

```
/Users/charliecarrera/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:1: DeprecationWarning:
.ix is deprecated. Please use
```

```
.loc for label based indexing or  
.iloc for positional indexing
```

See the documentation here:

```
http://pandas.pydata.org/pandas-docs/stable/indexing.html#ix-indexer-is-deprecated  
"""Entry point for launching an IPython kernel.
```

```
Out[392]: 5      1  
          5      1  
          Name: MasVnrType, dtype: int64
```

```
In [393]: data1['MasVnrArea'][data1['MasVnrArea'].isnull()] = [data1.groupby('MasVnrType')['MasVnrArea'].mean()  
/Users/charliecarrera/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:1: SettingWithCopyError:  
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#ix-indexer-is-deprecated>
"""Entry point for launching an IPython kernel.

```
In [394]: data1['LotFrontage'][data1['LotFrontage'].isnull()] = [data1.groupby('LotShape')['LotFrontage'].mean()  
/Users/charliecarrera/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:1: SettingWithCopyError:  
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#ix-indexer-is-deprecated>
"""Entry point for launching an IPython kernel.

```
In [ ]: data1
```

```
In [396]: train_kaggle = data1[data1['SalePrice'].notnull()]
```

```
In [398]: test_kaggle = data1[data1['SalePrice'].isnull()]
```

```
In [400]: test_train = train_kaggle.drop('SalePrice', axis = 1)
```

```
In [ ]: #Multivariate regression is used for the predictions
```

```
In [413]: from sklearn.linear_model import LinearRegression  
          L = LinearRegression()
```

```
In [414]: L.fit(train_kaggle.drop('SalePrice', axis = 1), train_kaggle['SalePrice'])
```

```
Out[414]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

```
In [415]: from sklearn.metrics import r2_score, mean_squared_error
```

```
In [416]: r2_score(train_kaggle['SalePrice'], L.predict(train_kaggle.drop('SalePrice', axis = 1)))
```

```
Out[416]: 0.8477035183629901
```

```
In [426]: mean_squared_error(train_kaggle['SalePrice'], L.predict(train_kaggle.drop('SalePrice
```

```
Out[426]: 30991.962054331976
```

```
In [418]: predictions = L.predict(test_kaggle)
```

```
In [419]: len(predictions)
```

```
Out[419]: 1459
```

```
In [420]: result = pd.DataFrame(columns = ['Id', 'SalePrice'])  
          result['Id'] = test_kaggle['Id']  
          result['SalePrice'] = predictions
```

```
Out[420]:
```

	Id	SalePrice
0	1461	103466.042422
1	1462	167918.528337
2	1463	177304.286094
3	1464	203285.743053
4	1465	186145.314042

```
In [422]: result.head()
```

```
Out[422]:
```

	Id	SalePrice
0	1461	103466.042422
1	1462	167918.528337
2	1463	177304.286094
3	1464	203285.743053
4	1465	186145.314042

```
In [ ]: #Convert result data to csv
```

```
In [424]: result.to_csv('SalePrice Kaggle.csv', index = False)
```